

## ▼ Declare feature vector and target variable

```
X = cc_df.drop(['Biopsy'], axis=1)
y = cc_df['Biopsy']
```

## ▼ Split data into separate training and test set

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 0)

X_train.shape, X_test.shape

((668, 35), (167, 35))

y_train.shape

(668,)
```

## ▼ Engineering outliers

```
def max_value(df3, variable, top):
    return np.where(df3[variable]>top, top, df3[variable])

for df3 in [X_train, X_test]:
    df3['Age'] = max_value(df3, 'Age', 48.5)
    df3['Number_of_sexual_partners'] = max_value(df3, 'Number_of_sexual_partners', 44.5)
    df3['First_sexual_intercourse'] = max_value(df3, 'First_sexual_intercourse', 22.5)
    df3['Num_of_pregnancies'] = max_value(df3, 'Num_of_pregnancies', 6)

X_train.Age.max(), X_test.Age.max()

(48.5, 48.5)

X_train.Number_of_sexual_partners.max(), X_test.Number_of_sexual_partners.max()

(28.0, 5.0)

X_train.First_sexual_intercourse.max(), X_test.First_sexual_intercourse.max()

(22.5, 22.5)

X_train.Num_of_pregnancies.max(), X_test.Num_of_pregnancies.max()

(6.0, 6.0)
```

## ▼ Feature Scaling

```
X_train.describe()
```

8 rows x 35 columns

```
from sklearn.preprocessing import MinMaxScaler
```

- Model Training

[illegible]

```
logreg= LogisticRegression(solver='liblinear', random_state=0)
```

```
LogisticRegression
LogisticRegression(random_state=0, solver='liblinear')
```

- ✓ Predict results

y\_pred\_test

```
0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0,
0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])
```

```
logreg.predict_proba(X_test)[:,0] # probabilities of getting no biopsy
```

```
array([0.91992733, 0.98901255, 0.99213392, 0.99338499, 0.96350935,
0.98492135, 0.68396716, 0.99280932, 0.98619543, 0.9800177 ,
0.96957308, 0.9884378 , 0.98259612, 0.9885742 , 0.46033742,
0.9812865 , 0.99320811, 0.98325562, 0.98527004, 0.97930573,
0.97639997, 0.98044322, 0.91443637, 0.98776232, 0.98894744,
0.95692314, 0.98609478, 0.98749521, 0.99896489, 0.27489498,
0.96384501, 0.98325726, 0.88303876, 0.99172243, 0.93884554,
0.98294622, 0.98942283, 0.98886892, 0.99482286, 0.98490239,
0.98256677, 0.98364479, 0.9813101 , 0.98599757, 0.99631075,
0.9879477 , 0.52490908, 0.9746085 , 0.9904325 , 0.98865874,
0.98557267, 0.98447354, 0.94496033, 0.98596821, 0.99735405,
0.91022135, 0.97966427, 0.98364959, 0.9874529 , 0.9843784 ,
0.98908174, 0.96781357, 0.99008667, 0.98170679, 0.98179538,
0.99195909, 0.97511538, 0.96864293, 0.98391754, 0.98073402,
0.46809188, 0.98836638, 0.9768906 , 0.97928647, 0.95045374,
0.9878481 , 0.98824817, 0.57344977, 0.98684004, 0.9918064 ,
0.98865707, 0.45668094, 0.98791917, 0.9893904 , 0.99117006,
0.97200749, 0.98954389, 0.98306286, 0.98920591, 0.98310232,
0.98867978, 0.97865247, 0.97546575, 0.98849014, 0.98212838,
0.98356008, 0.98980109, 0.98849196, 0.99568845, 0.98961769,
0.21202714, 0.9654617 , 0.99136161, 0.99135604, 0.98278209,
0.98105385, 0.415899 , 0.99049651, 0.98217583, 0.97771181,
0.98954869, 0.99218039, 0.97364164, 0.98029408, 0.27568103,
0.98888715, 0.25541798, 0.97489767, 0.98355687, 0.99334744,
0.99244108, 0.99356519, 0.98292325, 0.98145791, 0.98705342,
0.98095867, 0.98549552, 0.98484295, 0.99117503, 0.98576044,
0.98715349, 0.42116867, 0.99606284, 0.98257991, 0.99234688,
0.98484498, 0.26332961, 0.98598453, 0.99339527, 0.98658895,
0.98471595, 0.98409526, 0.99176459, 0.99494561, 0.99692705,
0.98928807, 0.98723822, 0.99027939, 0.99476763, 0.98940565,
0.98304047, 0.96238851, 0.99265516, 0.99630946, 0.98989316,
0.99517441, 0.98481641, 0.51253167, 0.99027703, 0.60745069,
0.98894497, 0.97933178, 0.98356808, 0.98392162, 0.97376315,
0.9851037 , 0.9902956 ])
```

```
logreg.predict_proba(X_test)[:,1] # probabilities of getting biopsy
```

```
array([0.08007267, 0.01098745, 0.00786608, 0.00661501, 0.03649065,
0.01507865, 0.31603284, 0.00719068, 0.01380457, 0.0199823 ,
0.03042692, 0.0115622 , 0.01740388, 0.0114258 , 0.53966258,
0.0187135 , 0.00679189, 0.01674438, 0.01472996, 0.02069427,
0.02360003, 0.01955678, 0.08556363, 0.01223768, 0.01105256,
0.04307686, 0.01390522, 0.01250479, 0.00103511, 0.72510502,
0.03615499, 0.01674274, 0.11696124, 0.00827757, 0.06115446,
0.01705378, 0.01057717, 0.01113108, 0.00517714, 0.01509761,
0.01743323, 0.01635521, 0.0186899 , 0.01400243, 0.00368925,
0.0120523 , 0.47509092, 0.0253915 , 0.0095675 , 0.01134126,
0.01442733, 0.01552646, 0.05503967, 0.01403179, 0.00264595,
0.08977865, 0.02033573, 0.01635041, 0.0125471 , 0.0156216 ,
0.01091826, 0.03218643, 0.00991333, 0.01829321, 0.01820462,
0.00804091, 0.02488462, 0.03135707, 0.01608246, 0.01926598,
0.53190812, 0.01163362, 0.0231094 , 0.02071353, 0.04954626,
0.0121519 , 0.01175183, 0.42655023, 0.01315996, 0.0081936 ,
0.01134293, 0.54331906, 0.01208083, 0.0106096 , 0.00882994,
0.02799251, 0.01045611, 0.01693714, 0.01079409, 0.01689768,
0.01132022, 0.02134753, 0.02453425, 0.01150986, 0.01787162,
0.01643992, 0.01019891, 0.01150804, 0.00431155, 0.01038231,
0.78797286, 0.0345383 , 0.00863839, 0.00864396, 0.01721791,
0.01894615, 0.584101 , 0.00950349, 0.01782417, 0.02228819,
0.01045131, 0.00781961, 0.02635836, 0.01970592, 0.72431897,
0.01111285, 0.74458202, 0.02510233, 0.01644313, 0.00665256,
0.00755892, 0.00643481, 0.01707675, 0.01854209, 0.01294658,
0.01904133, 0.01450448, 0.01515705, 0.00882497, 0.01423956,
0.01284651, 0.57883133, 0.00393716, 0.01742009, 0.00765312,
0.01515502, 0.73667039, 0.01401547, 0.00660473, 0.01341105,
0.01528405, 0.01590474, 0.00823541, 0.00505439, 0.00307295,
0.01071193, 0.01276178, 0.00972061, 0.00523237, 0.01059435,
0.01695953, 0.03761149, 0.00734484, 0.00369054, 0.01010684,
0.00482559, 0.01518359, 0.48746833, 0.00972297, 0.39254931,
0.01105503, 0.02066822, 0.01643192, 0.01607838, 0.02623685,
0.0148963 , 0.0097044 ])
```

- ✓ Check accuracy score

```
from sklearn.metrics import accuracy_score

print('Model accuracy score: {0:0.4f}'.format(accuracy_score(y_test, y_pred_test)))

Model accuracy score: 0.9401
```

- ✓ Compare the train-set and test-set accuracy

[illegible]

- ✓ Check for overfitting and underfitting

```
print('Training set score: {:.4f}'.format(logreg.score(X_train, y_train)))

print('Test set score: {:.4f}'.format(logreg.score(X_test, y_test)))

Training set score: 0.9671
Test set score: 0.9401

logreg100 = LogisticRegression(C=100, solver='liblinear', random_state=0)

logreg100.fit(X_train, y_train)



LogisticRegression



LogisticRegression(C=100, random_state=0, solver='liblinear')



print('Training set score: {:.4f}'.format(logreg100.score(X_train, y_train)))

print('Test set score: {:.4f}'.format(logreg100.score(X_test, y_test)))
```

```
Training set score: 0.9746
Test set score: 0.9281
```

```
logreg001 = LogisticRegression(C=0.01, solver='liblinear', random_state=0)
logreg001.fit(X_train, y_train)
```

```
▼ LogisticRegression
LogisticRegression(C=0.01, random_state=0, solver='liblinear')
```

```
print('Training set score: {:.4f}'.format(logreg001.score(X_train, y_train)))
print('Test set score: {:.4f}'.format(logreg001.score(X_test, y_test)))

Training set score: 0.9341
Test set score: 0.9401
```

## ▼ Compare model accuracy with null accuracy

```
y_test.value_counts()

Biopsy
0    157
1     10
Name: count, dtype: int64

null_accuracy = (157/(157+10))

print('Null accuracy score: {0:.4f}'.format(null_accuracy))

Null accuracy score: 0.9401
```

```
from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test, y_pred_test)

print('Confusion matrix\n\n', cm)

print('\nTrue Positives(TP) = ', cm[0,0])

print('\nTrue Negatives(TN) = ', cm[1,1])

print('\nFalse Positives(FP) = ', cm[0,1])

print('\nFalse Negatives(FN) = ', cm[1,0])
```

```
Confusion matrix

[[152   5]
 [  5   5]

True Positives(TP) = 152

True Negatives(TN) =  5

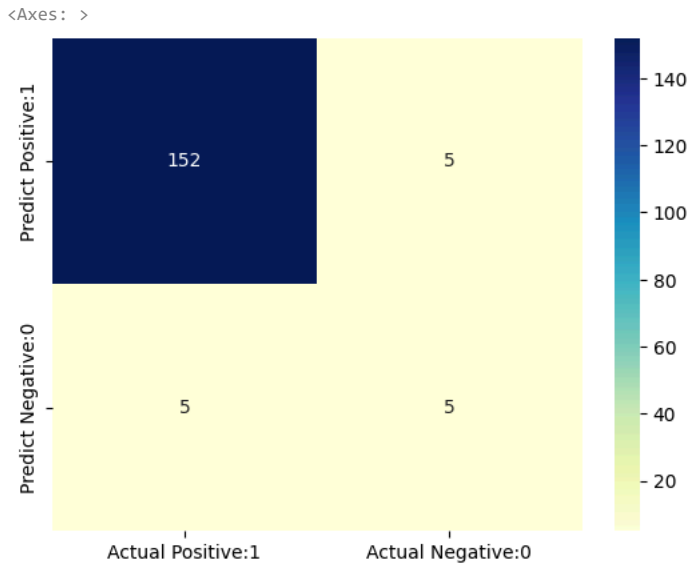
False Positives(FP) =  5

False Negatives(FN) =  5
```

▼ After creating confusion matrix, it could be seen that out of 167 results, there are only 5 results each for Type 1 and 2 errors.

```
cm_matrix = pd.DataFrame(data=cm, columns=['Actual Positive:1', 'Actual Negative:0'],
                        index=['Predict Positive:1', 'Predict Negative:0'])

sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='YlGnBu')
```



## Classification Report

```
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred_test))
```

	precision	recall	f1-score	support
0	0.97	0.97	0.97	157
1	0.50	0.50	0.50	10
accuracy			0.94	167
macro avg	0.73	0.73	0.73	167
weighted avg	0.94	0.94	0.94	167

## Classification Accuracy

```
TP = cm[0,0]
TN = cm[1,1]
FP = cm[0,1]
FN = cm[1,0]

classification_accuracy = (TP + TN) / float(TP + TN + FP + FN)

print('Classification accuracy : {0:0.4f}'.format(classification_accuracy))

Classification accuracy : 0.9401
```

## Classification Error

```
classification_error = (FP + FN) / float(TP + TN + FP + FN)

print('Classification error : {0:0.4f}'.format(classification_error))

Classification error : 0.0599
```

## Precision

```
precision = TP / float(TP+FP)

print('Precision: {0:0.4f}'.format(precision))
```

Precision: 0.9682

## Recall

```
recall = TP / float(TP + FN)

print('Recall or Sensitivity : {0:0.4f}'.format(recall))
```

Recall or Sensitivity : 0.9682

## True Positive Rate

```
true_positive_rate = TP / float(TP + FN)
print('True Positive Rate : {0:0.4f}'.format(true_positive_rate))
```

True Positive Rate : 0.9682

## Specificity

```
specificity = TN / (TN+FP)
print('Specifity : {0:0.4f}'.format(specificity))
```

Specifity : 0.5000

## Adjusting the threshold level

```
y_pred_prob = logreg.predict_proba(X_test)[0:10]
```

```
y_pred_prob
```

```
array([[0.91992733, 0.08007267],
       [0.98901255, 0.01098745],
       [0.99213392, 0.00786608],
       [0.99338499, 0.00661501],
       [0.96350935, 0.03649065],
       [0.98492135, 0.01507865],
       [0.68396716, 0.31603284],
       [0.99280932, 0.00719068],
       [0.98619543, 0.01380457],
       [0.9800177 , 0.0199823 ]])
```

```
y_pred_prob_df = pd.DataFrame(data=y_pred_prob, columns=['Probability of no Biopsy', 'Probability of a Biopsy'])
```

```
y_pred_prob_df
```

	Probability of no Biopsy	Probability of a Biopsy
0	0.919927	0.080073
1	0.989013	0.010987
2	0.992134	0.007866
3	0.993385	0.006615
4	0.963509	0.036491
5	0.984921	0.015079
6	0.683967	0.316033
7	0.992809	0.007191
8	0.986195	0.013805
9	0.980018	0.019982

Next steps: [View recommended plots](#)

```
logreg.predict_proba(X_test)[0:10, 1]

array([0.08007267, 0.01098745, 0.00786608, 0.00661501, 0.03649065,
       0.01507865, 0.31603284, 0.00719068, 0.01380457, 0.0199823 ])

y_pred1 = logreg.predict_proba(X_test)[: , 1]

from sklearn.preprocessing import binarize

for i in range(1,5):
    cm1=0

    y_pred1 = logreg.predict_proba(X_test)[: ,1]

    y_pred1 = y_pred1.reshape(-1,1)

    y_pred2 = binarize(y_pred1, threshold=(i/10))

    cm1 = confusion_matrix(y_test, y_pred2)
    print ('With',i/10, 'threshold the Confusion Matrix is ', '\n\n', cm1, '\n\n',
          'with', cm1[0,0]+cm1[1,1], 'correct predictions, ', '\n\n',
          cm1[0,1], 'Type I errors( False Positives), ', '\n\n',
          cm1[1,0], 'Type II errors( False Negatives), ', '\n\n',
          'Accuracy score: ', (accuracy_score(y_test, y_pred2)), '\n\n',
          'Sensitivity: ', cm1[1,1]/(float(cm1[1,1]+cm1[1,0])), '\n\n',
          'Specificity: ', cm1[0,0]/(float(cm1[0,0]+cm1[0,1])), '\n\n',
          '=====', '\n\n')

    [[149   8]
     [  3   7]]

    with 156 correct predictions,

    8 Type I errors( False Positives),

    3 Type II errors( False Negatives),

    Accuracy score:  0.9341317365269461

    Sensitivity:  0.7

    Specificity:  0.9490445859872612

    =====

With 0.3 threshold the Confusion Matrix is

[[149   8]
 [  3   7]]

with 156 correct predictions,

8 Type I errors( False Positives),
```



Accuracy score: 0.9341317365269461

Sensitivity: 0.6

Specificity: 0.9554140127388535

=====

```
plt.rcParams['font.size'] = 12
```

```
plt.hist(y_pred1, bins = 10)
```

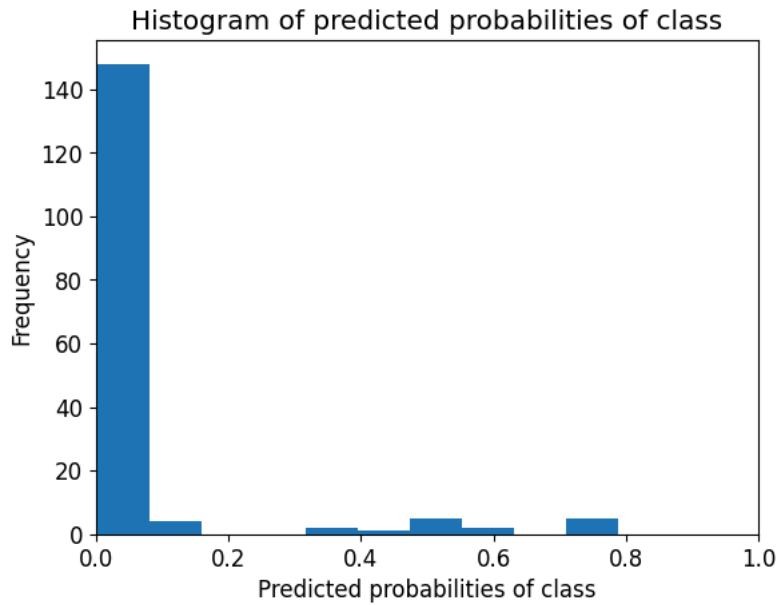
```
plt.title('Histogram of predicted probabilities of class')
```

```
plt.xlim(0,1)
```

```
plt.xlabel('Predicted probabilities of class')
```

```
plt.ylabel('Frequency')
```

```
Text(0, 0.5, 'Frequency')
```



## ✓ ROC Curve

```
from sklearn.metrics import roc_curve
```

```
fpr, tpr, thresholds = roc_curve(y_test, y_pred1, pos_label = 'Yes')
```

```
plt.figure(figsize=(6,4))
```

```
plt.plot(fpr,tpr, linewidth=2)
```

```
plt.plot([0,1], [0,1], 'k--')
```


```
plt.rcParams['font.size'] = 12
```

```
plt.title('ROC Curve for Biopsy classifier')
```

```
plt.xlabel('False Positive Rate (1 - Specificity)')
```

```
plt.ylabel('True Positive Rate (Sensitivity)')
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_ranking.py:1029: UndefinedMetricWarning: No positive samples in y_true, true pc
warnings.warn(
Text(0, 0.5, 'True Positive Rate (Sensitivity)')
```



No positive samples in  $y_{\text{true}}$  so there is no curve.

## ✓ ROC-AUC

```
from sklearn.metrics import roc_auc_score

ROC_AUC = roc_auc_score(y_test, y_pred1)

print('ROC AUC : {:.4f}'.format(ROC_AUC))

ROC AUC : 0.9153
```

## ✓ Cross validated ROC-AUC

```
from sklearn.model_selection import cross_val_score

Cross_validated_ROC_AUC = cross_val_score(logreg, X_train, y_train, cv=5, scoring='roc_auc').mean()

print('Cross validated ROC AUC : {:.4f}'.format(Cross_validated_ROC_AUC))

Cross validated ROC AUC : 0.9434
```

## ✓ k-Fold Cross Validation

```
from sklearn.model_selection import cross_val_score

scores = cross_val_score(logreg, X_train, y_train, cv=5, scoring = 'accuracy')

print('Cross-validation scores:{}'.format(scores))

Cross-validation scores:[0.95522388 0.94776119 0.94776119 0.96240602 0.94736842]

print('Average cross-validation score: {:.4f}'.format(scores.mean()))

Average cross-validation score: 0.9521
```

## ✓ Hyperparameter Optimization using GridSearch CV

```

from sklearn.model_selection import GridSearchCV

parameters = [{'penalty':['l1','l2']],
               {'C':[1, 10, 100, 1000]]}

grid_search = GridSearchCV(estimator = logreg,
                           param_grid = parameters,
                           scoring = 'accuracy',
                           cv = 5,
                           verbose=0)

grid_search.fit(X_train, y_train)

/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:378: FitFailedWarning:
10 fits failed out of a total of 30.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:
-----
5 fits failed with the following error:
Traceback (most recent call last):
  File "/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1160, in fit
    self._validate_params()
  File "/usr/local/lib/python3.10/dist-packages/sklearn/base.py", line 600, in _validate_params
    validate_parameter_constraints(
  File "/usr/local/lib/python3.10/dist-packages/sklearn/utils/_param_validation.py", line 97, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'penalty' parameter of LogisticRegression must be a str among {'l1', 'elastic

-----
5 fits failed with the following error:
Traceback (most recent call last):
  File "/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1160, in fit
    self._validate_params()
  File "/usr/local/lib/python3.10/dist-packages/sklearn/base.py", line 600, in _validate_params
    validate_parameter_constraints(
  File "/usr/local/lib/python3.10/dist-packages/sklearn/utils/_param_validation.py", line 97, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'penalty' parameter of LogisticRegression must be a str among {'l1', 'elastic

warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:952: UserWarning: One or more of the test scores are non-fini
warnings.warn(

```



```

print('GridSearch CV best score : {:.4f}\n\n'.format(grid_search.best_score_))

print('Parameters that give the best results :', '\n\n', (grid_search.best_params_))

print('\n\nEstimator that was chosen by the search :', '\n\n', (grid_search.best_estimator_))

GridSearch CV best score : 0.9551

Parameters that give the best results :

{'C': 10}

Estimator that was chosen by the search :

LogisticRegression(C=10, random_state=0, solver='liblinear')

print('GridSearch CV score on test set: {0:0.4f}'.format(grid_search.score(X_test, y_test)))

GridSearch CV score on test set: 0.9281

```