

✓ Name: Dylan James N. Dejoras

Section: CPE22S3

source:
<https://archive.ics.uci.edu/dataset/10/automobile>
<https://archive.ics.uci.edu/dataset/109/wine>

✓ Setup

```
pip install ucimlrepo

Requirement already satisfied: ucimlrepo in /usr/local/lib/python3.10/dist-packages (0.0.6)

import pandas as pd
import numpy as np

from ucimlrepo import fetch_ucirepo

# fetch dataset
automobile = fetch_ucirepo(id=10)

# data (as pandas dataframes)
X = automobile.data.features
y = automobile.data.targets

# metadata
print(automobile.metadata)

# variable information
print(automobile.variables)
```

	{'uci_id': 10, 'name': 'Automobile', 'repository_url': 'https://archive.ics.uci.edu/dataset/10/automobile', 'data_url': 'https://archive
	name role type demographic \
0	price Feature Continuous None
1	highway-mpg Feature Continuous None
2	city-mpg Feature Continuous None
3	peak-rpm Feature Continuous None
4	horsepower Feature Continuous None
5	compression-ratio Feature Continuous None
6	stroke Feature Continuous None
7	bore Feature Continuous None
8	fuel-system Feature Categorical None
9	engine-size Feature Continuous None
10	num-of-cylinders Feature Integer None
11	engine-type Feature Categorical None
12	curb-weight Feature Continuous None
13	height Feature Continuous None
14	width Feature Continuous None
15	length Feature Continuous None
16	wheel-base Feature Continuous None
17	engine-location Feature Binary None
18	drive-wheels Feature Categorical None
19	body-style Feature Categorical None
20	num-of-doors Feature Integer None
21	aspiration Feature Binary None
22	fuel-type Feature Binary None
23	make Feature Categorical None
24	normalized-losses Feature Continuous None
25	symboling Target Integer None

	description	units	missing_values
0	continuous from 5118 to 45400	None	yes
1	continuous from 16 to 54	None	no
2	continuous from 13 to 49	None	no
3	continuous from 4150 to 6600	None	yes
4	continuous from 48 to 288	None	yes
5	continuous from 7 to 23	None	no
6	continuous from 2.07 to 4.17	None	yes
7	continuous from 2.54 to 3.94	None	yes

```

8      1bbl, 2bbl, 4bbl, idi, mfi, mpfi, spdi, spfi None      no
9      continuous from 61 to 326 None      no
10     eight, five, four, six, three, twelve, two None      no
11     dohc, dohcv, l, ohc, ohcf, ohcv, rotor None      no
12     continuous from 1488 to 4066 None      no
13     continuous from 47.8 to 59.8 None      no
14     continuous from 60.3 to 72.3 None      no
15     continuous from 141.1 to 208.1 None      no
16     continuous from 86.6 120.9 None      no
17     front, rear None      no
18     4wd, fwd, rwd None      no
19     hardtop, wagon, sedan, hatchback, convertible None      no
20     four, two None      yes
21     std, turbo None      no
22     diesel, gas None      no
23     alfa-romero, audi, bmw, chevrolet, dodge, hond... None      no
24     continuous from 65 to 256 None      yes
25     -3, -2, -1, 0, 1, 2, 3 None      no

```

```

dataFrames = [X,y]
df = pd.concat(dataFrames, axis = 1)
df

```

	price	highway- mpg	city- mpg	peak- rpm	horsepower	compression- ratio	stroke	bore	fuel- system	en
0	13495.0	27	21	5000.0	111.0	9.0	2.68	3.47	mpfi	
1	16500.0	27	21	5000.0	111.0	9.0	2.68	3.47	mpfi	
2	16500.0	26	19	5000.0	154.0	9.0	3.47	2.68	mpfi	
3	13950.0	30	24	5500.0	102.0	10.0	3.40	3.19	mpfi	
4	17450.0	22	18	5500.0	115.0	8.0	3.40	3.19	mpfi	
...	
200	16845.0	28	23	5400.0	114.0	9.5	3.15	3.78	mpfi	
201	19045.0	25	19	5300.0	160.0	8.7	3.15	3.78	mpfi	
202	21485.0	23	18	5500.0	134.0	8.8	2.87	3.58	mpfi	
203	22470.0	27	26	4800.0	106.0	23.0	3.40	3.01	idi	
204	22625.0	25	19	5400.0	114.0	9.5	3.15	3.78	mpfi	

205 rows × 26 columns

```
auto_df = df.copy()
```

```

def check_duplicates(df):
    if df[df.duplicated()].shape[0] != 0:
        print(df[df.duplicated()].shape[0])
    else:
        print("No existing duplicates")
check_duplicates(auto_df)

```

No existing duplicates

```
df.isnull().sum()
```

```

price      4
highway-mpg 0
city-mpg   0
peak-rpm   2
horsepower 2
compression-ratio 0
stroke     4
bore       4
fuel-system 0
engine-size 0
num-of-cylinders 0
engine-type 0
curb-weight 0

```

```

height          0
width           0
length          0
wheel-base     0
engine-location 0
drive-wheels    0
body-style      0
num-of-doors    2
aspiration      0
fuel-type       0
make            0
normalized-losses 41
symboling       0
dtype: int64

```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   price                 201 non-null   float64
1   highway-mpg           205 non-null   int64
2   city-mpg              205 non-null   int64
3   peak-rpm              203 non-null   float64
4   horsepower            203 non-null   float64
5   compression-ratio     205 non-null   float64
6   stroke                201 non-null   float64
7   bore                  201 non-null   float64
8   fuel-system           205 non-null   object
9   engine-size           205 non-null   int64
10  num-of-cylinders      205 non-null   int64
11  engine-type           205 non-null   object
12  curb-weight           205 non-null   int64
13  height                205 non-null   float64
14  width                 205 non-null   float64
15  length                205 non-null   float64
16  wheel-base           205 non-null   float64
17  engine-location       205 non-null   object
18  drive-wheels          205 non-null   object
19  body-style            205 non-null   object
20  num-of-doors          203 non-null   float64
21  aspiration            205 non-null   object
22  fuel-type             205 non-null   object
23  make                  205 non-null   object
24  normalized-losses     164 non-null   float64
25  symboling             205 non-null   int64
dtypes: float64(12), int64(6), object(8)
memory usage: 41.8+ KB

```

```

def check_value_counts(column):
    print(df.value_counts(column))

```

```
check_value_counts('normalized-losses')
```

```

normalized-losses
161.0    11
91.0     8
150.0    7
128.0    6
104.0    6
134.0    6
95.0     5
94.0     5
74.0     5
65.0     5
103.0    5
85.0     5
168.0    5
102.0    5
122.0    4
148.0    4
106.0    4
118.0    4
93.0     4
101.0    3
154.0    3
115.0    3
83.0     3
125.0    3

```

```

137.0    3
87.0     2
188.0    2
158.0    2
153.0    2
81.0     2
145.0    2
192.0    2
89.0     2
129.0    2
194.0    2
197.0    2
119.0    2
113.0    2
110.0    2
108.0    2
164.0    2
186.0    1
231.0    1
142.0    1
77.0     1
78.0     1
98.0     1
90.0     1
121.0    1
107.0    1
256.0    1
Name: count, dtype: int64

```

```

def fill_missing_values(df, columns):
    for col in columns:
        df[col] = df[col].fillna(df[col].median())
    return df

```

```
na_counts = df.isnull().sum()
```

```
columns_with_na = na_counts[na_counts > 0].index.tolist()
```

```
columns_with_na
```

```

['price',
 'peak-rpm',
 'horsepower',
 'stroke',
 'bore',
 'num-of-doors',
 'normalized-losses']

```

```
auto_df = fill_missing_values(df, columns_with_na)
```

```
auto_df.isnull().sum()
```

```

price           0
highway-mpg     0
city-mpg        0
peak-rpm        0
horsepower      0
compression-ratio 0
stroke          0
bore            0
fuel-system     0
engine-size     0
num-of-cylinders 0
engine-type     0
curb-weight     0
height          0
width           0
length          0
wheel-base     0
engine-location 0
drive-wheels    0
body-style     0
num-of-doors    0
aspiration     0
fuel-type       0
make           0
normalized-losses 0

```

```

symboling      0
dtype: int64

```

```
auto_df.dtypes
```

```

price          float64
highway-mpg    int64
city-mpg       int64
peak-rpm       float64
horsepower     float64
compression-ratio float64
stroke         float64
bore           float64
fuel-system    object
engine-size    int64
num-of-cylinders int64
engine-type    object
curb-weight    int64
height         float64
width          float64
length         float64
wheel-base    float64
engine-location object
drive-wheels   object
body-style     object
num-of-doors   float64
aspiration     object
fuel-type      object
make           object
normalized-losses float64
symboling      int64
dtype: object

```

```

Fuel_system_type = list(df['fuel-system'].unique())
Engine_type = list(df['engine-type'].unique())
Engine_location_type = list(df['engine-location'].unique())
Drive_wheels_type = list(df['drive-wheels'].unique())
Body_style_type = list(df['body-style'].unique())
Aspiration_type = list(df['aspiration'].unique())
Fuel_type = list(df['fuel-type'].unique())
Make_type = list(df['make'].unique())

```

```

df['fuel-system'] = df.apply(lambda x: Fuel_system_type.index(x['fuel-system']) + 1, axis=1)
df['engine-type'] = df.apply(lambda x: Engine_type.index(x['engine-type']) + 1, axis=1)
df['engine-location'] = df.apply(lambda x: Engine_location_type.index(x['engine-location']) + 1, axis=1)
df['drive-wheels'] = df.apply(lambda x: Drive_wheels_type.index(x['drive-wheels']) + 1, axis=1)
df['body-style'] = df.apply(lambda x: Body_style_type.index(x['body-style']) + 1, axis=1)
df['aspiration'] = df.apply(lambda x: Aspiration_type.index(x['aspiration']) + 1, axis=1)
df['fuel-type'] = df.apply(lambda x: Fuel_type.index(x['fuel-type']) + 1, axis=1)
df['make'] = df.apply(lambda x: Make_type.index(x['make']) + 1, axis=1)

```

```
auto_df.dtypes
```

```


price          float64
highway-mpg    int64
city-mpg       int64
peak-rpm       float64
horsepower     float64
compression-ratio float64
stroke         float64
bore           float64
fuel-system    int64
engine-size    int64
num-of-cylinders int64
engine-type    int64
curb-weight    int64
height         float64
width          float64
length         float64
wheel-base    float64
engine-location int64
drive-wheels   int64
body-style     int64
num-of-doors   float64
aspiration     int64
fuel-type      int64
make           int64
normalized-losses float64

```

```
symboling      int64  
dtype: object
```

```
%matplotlib inline  
import seaborn as sns  
import matplotlib.pyplot as plt  
  
plt.figure(figsize=(20, 20))  
sns.heatmap(auto_df.corr(), annot=True)
```

<Axes: >



```

X = df.drop('horsepower', axis=1)
y = df['horsepower']

print("X=",X.shape,"\ny=",y.shape)

X= (205, 25)
y= (205,)

from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.linear_model import LinearRegression

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=101)

X_train.shape

(143, 25)

X_test.shape

(62, 25)

model = LinearRegression()

model.fit(X_train, y_train)

▼ LinearRegression
LinearRegression()

model.coef_

array([-4.93081843e-04,  3.50592614e+00, -4.34114068e+00,  1.96062716e-02,
        5.22089428e+00,  7.30210367e+00,  2.44507227e+01, -7.58307630e-01,
        2.63253993e-01,  1.02400233e+01,  1.45818418e+00,  1.84326982e-02,
       -3.67802358e-02,  1.54274355e+00, -7.75070944e-02, -8.05226840e-01,
        2.00206382e+01, -9.60977760e+00, -9.58622509e-01, -1.46993056e+00,
        3.40275905e+01, -9.03368882e+01,  6.30398597e-03, -4.63055804e-02,
       -2.36585941e-01])

pd.DataFrame(model.coef_, X.columns, columns=['Coefficients'])

```


Coefficients	
price	-0.000493
highway-mpg	3.505926
city-mpg	-4.341141
peak-rpm	0.019606
compression-ratio	5.220894
stroke	7.302104
bore	24.450723
fuel-system	-0.758308
engine-size	0.263254
num-of-cylinders	10.240023
engine-type	1.458184
curb-weight	0.018433
height	-0.036780
width	1.542744
length	-0.077507
wheel-base	-0.805227
engine-location	20.020638
drive-wheels	-9.609778
body-style	-0.958623
num-of-doors	-1.469931
aspiration	34.027590
fuel-type	-90.336888
make	0.006304
normalized-losses	-0.046306
symboling	-0.236586

```
y_pred = model.predict(X_test)

MAE = metrics.mean_absolute_error(y_test, y_pred)
MSE = metrics.mean_squared_error(y_test, y_pred)
RMSE = np.sqrt(MSE)
```

Average Error

```
MAE

9.592603226136523
```

Mean Square Error

```
MSE

180.35954074625744
```

Root Square Mean Error

```
RMSE

13.429800473062041
```

▼ average of horsepower

```
df['horsepower'].mean()

104.16585365853659
```

▼ subtracting y test and y pred yields test residual

```
test_residual = y_test - y_pred
```

```
!pip install hvplot
```

```
Requirement already satisfied: hvplot in /usr/local/lib/python3.10/dist-packages (0.9.2)
Requirement already satisfied: bokeh>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from hvplot) (3.3.4)
Requirement already satisfied: colorcet>=2 in /usr/local/lib/python3.10/dist-packages (from hvplot) (3.1.0)
Requirement already satisfied: holoviews>=1.11.0 in /usr/local/lib/python3.10/dist-packages (from hvplot) (1.17.1)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from hvplot) (2.0.3)
Requirement already satisfied: numpy>=1.15 in /usr/local/lib/python3.10/dist-packages (from hvplot) (1.25.2)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from hvplot) (24.0)
Requirement already satisfied: panel>=0.11.0 in /usr/local/lib/python3.10/dist-packages (from hvplot) (1.3.8)
Requirement already satisfied: param<3.0,>=1.12.0 in /usr/local/lib/python3.10/dist-packages (from hvplot) (2.1.0)
Requirement already satisfied: Jinja2>=2.9 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (3.1.3)
Requirement already satisfied: contourpy>=1 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (1.2.1)
Requirement already satisfied: pillow>=7.1.0 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (9.4.0)
Requirement already satisfied: PyYAML>=3.10 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (6.0.1)
Requirement already satisfied: tornado>=5.1 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (6.3.3)
Requirement already satisfied: xyzservices>=2021.09.1 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (2024.4.0)
Requirement already satisfied: pyviz-comms>=0.7.4 in /usr/local/lib/python3.10/dist-packages (from holoviews>=1.11.0->hvplot) (3.0.2)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas->hvplot) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->hvplot) (2023.4)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas->hvplot) (2024.1)
Requirement already satisfied: markdown in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (3.6)
Requirement already satisfied: markdown-it-py in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (3.0.0)
Requirement already satisfied: linkify-it-py in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (2.0.3)
Requirement already satisfied: mdit-py-plugins in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (0.4.0)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (2.31.0)
Requirement already satisfied: tqdm>=4.48.0 in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (4.66.2)
Requirement already satisfied: bleach in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (6.1.0)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (4.11.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from Jinja2>=2.9->bokeh>=1.0.0->hvplot) (2.1.
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas->hvplot) (1.16.0)
Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-packages (from bleach->panel>=0.11.0->hvplot) (0.5.1)
Requirement already satisfied: uc-micro-py in /usr/local/lib/python3.10/dist-packages (from linkify-it-py->panel>=0.11.0->hvplot) (1.0.3)
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.10/dist-packages (from markdown-it-py->panel>=0.11.0->hvplot) (0.1.2)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->panel>=0.11.0->hvplot) (3.7)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->panel>=0.11.0->hvplot) (3.7)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests->panel>=0.11.0->hvplot) (2.0
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests->panel>=0.11.0->hvplot) (202
```


```
import hvplot.pandas
```

```
pd.DataFrame({'Error Values': (test_residual)}).hvplot.kde()
```

```
/usr/local/lib/python3.10/dist-packages/holoviews/core/util.py:1585: PanelDeprecationWarning: 'param_value_if_widget' is deprecated and
value = param_value_if_widget(value)
```


```
sns.displot(test_residual, bins=25, kde=True)
```

```
<seaborn.axisgrid.FacetGrid at 0x7e92c81ed570>
```



```
sns.scatterplot(x=y_test, y=test_residual)
plt.axhline(y=0, color='r', ls='--')
```

```
<matplotlib.lines.Line2D at 0x7e92c8054c10>
```



```
# fetch dataset
wine = fetch_ucirepo(id=109)
```

```
# data (as pandas dataframes)
x = wine.data.features
Y = wine.data.targets
```

```
# metadata
print(wine.metadata)
```

```
# variable information
print(wine.variables)
```

```
{'uci_id': 109, 'name': 'Wine', 'repository_url': 'https://archive.ics.uci.edu/dataset/109/wine', 'data_url': 'https://archive.ics.uci.edu/dataset/109/wine'}
0      name      role      type demographic \
1      class  Target  Categorical      None
2      Alcohol Feature  Continuous      None
3      Malicacid Feature  Continuous      None
```

3	Ash	Feature	Continuous	None
4	Alcalinity_of_ash	Feature	Continuous	None
5	Magnesium	Feature	Integer	None
6	Total_phenols	Feature	Continuous	None
7	Flavanoids	Feature	Continuous	None
8	Nonflavanoid_phenols	Feature	Continuous	None
9	Proanthocyanins	Feature	Continuous	None
10	Color_intensity	Feature	Continuous	None
11	Hue	Feature	Continuous	None
12	0D280_0D315_of_diluted_wines	Feature	Continuous	None
13	Proline	Feature	Integer	None

	description	units	missing_values
0	None	None	no
1	None	None	no
2	None	None	no
3	None	None	no
4	None	None	no
5	None	None	no
6	None	None	no
7	None	None	no
8	None	None	no
9	None	None	no
10	None	None	no
11	None	None	no
12	None	None	no
13	None	None	no

```
dataFrames = [x,Y]
wine_df = pd.concat(dataFrames, axis = 1)
wine_df
```

	Alcohol	Malicacid	Ash	Alcalinity_of_ash	Magnesium	Total_phenols	Flavanoids	Nonflavanoid_phenols	Proanthocyanins	Color_inte
0	14.23	1.71	2.43	15.6	127	2.80	3.06	0.28	2.29	
1	13.20	1.78	2.14	11.2	100	2.65	2.76	0.26	1.28	
2	13.16	2.36	2.67	18.6	101	2.80	3.24	0.30	2.81	
3	14.37	1.95	2.50	16.8	113	3.85	3.49	0.24	2.18	
4	13.24	2.59	2.87	21.0	118	2.80	2.69	0.39	1.82	
...
173	13.71	5.65	2.45	20.5	95	1.68	0.61	0.52	1.06	
174	13.40	3.91	2.48	23.0	102	1.80	0.75	0.43	1.41	
175	13.27	4.28	2.26	20.0	120	1.59	0.69	0.43	1.35	
176	13.17	2.59	2.37	20.0	120	1.65	0.68	0.53	1.46	
177	14.13	4.10	2.74	24.5	96	2.05	0.76	0.56	1.35	

178 rows x 14 columns

Next steps:

 View recommended plots

```
check_duplicates(wine_df)


No existing duplicates
```

```
wine_df.isnull().sum()

Alcohol      0
Malicacid    0
Ash          0
Alcalinity_of_ash  0
Magnesium    0
Total_phenols  0
Flavanoids   0
Nonflavanoid_phenols  0
Proanthocyanins  0
Color_intensity  0
Hue          0
0D280_0D315_of_diluted_wines  0
Proline      0
class       0
dtype: int64
```

```
plt.figure(figsize=(20, 20))  
sns.heatmap(wine_df.corr(), annot=True)
```

<Axes: >



```
wine_df.shape
```

```
(178, 14)
```

```
wine_df.head()
```

	Alcohol	Malicacid	Ash	Alcalinity_of_ash	Magnesium	Total_phenols	Flavanoids	Nonflavanoid_phenols	Proanthocyanins	Color_intens
0	14.23	1.71	2.43	15.6	127	2.80	3.06	0.28	2.29	!
1	13.20	1.78	2.14	11.2	100	2.65	2.76	0.26	1.28	!
2	13.16	2.36	2.67	18.6	101	2.80	3.24	0.30	2.81	!
3	14.37	1.95	2.50	16.8	113	3.85	3.49	0.24	2.18	!
4	13.24	2.59	2.87	21.0	118	2.80	2.69	0.39	1.82	!

Next steps: [View recommended plots](#)

```
wine_df.dtypes
```

```
Alcohol          float64
Malicacid        float64
Ash              float64
Alcalinity_of_ash float64
Magnesium        int64
Total_phenols    float64
Flavanoids       float64
Nonflavanoid_phenols float64
Proanthocyanins  float64
Color_intensity  float64
Hue              float64
0D280_0D315_of_diluted_wines float64
Proline          int64
class            int64
dtype: object
```

✓ We would deal with continuous variables since there is no categorical variables as the usual variable for logistic regression.

```
categorical = [var for var in wine_df.columns if wine_df[var].dtype=='O']
print(categorical)
```

```
[]
```

```
numerical = [var for var in wine_df.columns if wine_df[var].dtype!='O']
print(numerical)
```

```
['Alcohol', 'Malicacid', 'Ash', 'Alcalinity_of_ash', 'Magnesium', 'Total_phenols', 'Flavanoids', 'Nonflavanoid_phenols', 'Proanthocyanin']
```

```
wine_df[categorical].sum()
```

```
Series([], dtype: float64)
```

```
wine_df[numerical].isnull().sum()
```

```
Alcohol          0
Malicacid        0
Ash              0
Alcalinity_of_ash 0
```

```

Magnesium          0
Total_phenols      0
Flavanoids         0
Nonflavanoid_phenols 0
Proanthocyanins    0
Color_intensity    0
Hue                0
0D280_0D315_of_diluted_wines 0
Proline            0
class              0
dtype: int64

```

✓ Outliers in our dataset

```
print(round(wine_df[numerical].describe()),2)
```

```

      Alcohol  Malicacid   Ash  Alcalinity_of_ash  Magnesium  Total_phenols  \
count    178.0      178.0  178.0             178.0      178.0         178.0
mean      13.0         2.0   2.0              19.0      100.0          2.0
std       1.0         1.0   0.0               3.0       14.0          1.0
min      11.0         1.0   1.0              11.0       70.0          1.0
25%      12.0         2.0   2.0              17.0       88.0          2.0
50%      13.0         2.0   2.0              20.0       98.0          2.0
75%      14.0         3.0   3.0              22.0      107.0          3.0
max      15.0         6.0   3.0              30.0      162.0          4.0

      Flavanoids  Nonflavanoid_phenols  Proanthocyanins  Color_intensity  \
count      178.0                178.0             178.0             178.0
mean         2.0                  0.0                2.0                5.0
std          1.0                  0.0                1.0                2.0
min          0.0                  0.0                0.0                1.0
25%          1.0                  0.0                1.0                3.0
50%          2.0                  0.0                2.0                5.0
75%          3.0                  0.0                2.0                6.0
max          5.0                  1.0                4.0             13.0

      Hue  0D280_0D315_of_diluted_wines  Proline  class
count    178.0                178.0      178.0    178.0
mean      1.0                  3.0      747.0      2.0
std       0.0                  1.0     315.0      1.0
min       0.0                  1.0     278.0      1.0
25%       1.0                  2.0     500.0      1.0
50%       1.0                  3.0     674.0      2.0
75%       1.0                  3.0     985.0      3.0
max       2.0                  4.0    1680.0      3.0  2

```

```
plt.figure(figsize=(15,10))
```

```

plt.subplot(3, 3, 1)
fig = wine_df.boxplot(column = 'Malicacid')
fig.set_title('')
fig.set_ylabel('Malicacid')

```

```

plt.subplot(3, 3, 2)
fig = wine_df.boxplot(column = 'Alcalinity_of_ash')
fig.set_title('')
fig.set_ylabel('Alcalinity_of_ash')

```

```

plt.subplot(3, 3, 3)
fig = wine_df.boxplot(column = 'Magnesium')
fig.set_title('')
fig.set_ylabel('Magnesium')

```

```

plt.subplot(3, 3, 4)
fig = wine_df.boxplot(column = 'Proanthocyanins')
fig.set_title('')
fig.set_ylabel('Proanthocyanins')

```

```

plt.subplot(3, 3, 5)
fig = wine_df.boxplot(column = 'Color_intensity')
fig.set_title('')
fig.set_ylabel('Color_intensity')


```

```

plt.subplot(3, 3, 6)
fig = wine_df.boxplot(column = 'Hue')
fig.set_title('')
fig.set_ylabel('Hue')

```


Text(0, 0.5, 'Hue')



```
plt.figure(figsize=(15,10))
```

```
plt.subplot(3, 3, 1)
fig = wine_df.Malicacid.hist(bins=10)
fig.set_xlabel(' ')
fig.set_ylabel('')
```

```
plt.subplot(3, 3, 2)
fig = wine_df.Alcalinity_of_ash.hist(bins=10)
fig.set_xlabel('Alcalinity_of_ash')
fig.set_ylabel('')
```


```
plt.subplot(3, 3, 3)
fig = wine_df.Magnesium.hist(bins=10)
fig.set_xlabel('Magnesium')
fig.set_ylabel('')
```

```
plt.subplot(3, 3, 4)
fig = wine_df.Proanthocyanins.hist(bins=10)
fig.set_xlabel('Proanthocyanins')
fig.set_ylabel('')
```

```
plt.subplot(3, 3, 5)
fig = wine_df.Color_intensity.hist(bins=10)
fig.set_xlabel('Color_intensity')
fig.set_ylabel('')
```

```
plt.subplot(3, 3, 6)
fig = wine_df.Hue.hist(bins=10)
fig.set_xlabel('Hue')
fig.set_ylabel('')
```

Text(0, 0.5, '')



```

IQR = wine_df.Malicacid.quantile(0.75) - wine_df.Malicacid.quantile(0.25)
Lower_fence = wine_df.Malicacid.quantile(0.25) - (IQR * 3)
Upper_fence = wine_df.Malicacid.quantile(0.75) + (IQR * 3)
print('Malicacid outliers are values < {lowerboundary} or > {upperboundary}'.format(lowerboundary=Lower_fence, upperboundary=Upper_fence))

Malicacid outliers are values < -2.8374999999999995 or > 7.522499999999999

IQR = wine_df.Alcalinity_of_ash.quantile(0.75) - wine_df.Alcalinity_of_ash.quantile(0.25)
Lower_fence = wine_df.Alcalinity_of_ash.quantile(0.25) - (IQR * 3)
Upper_fence = wine_df.Alcalinity_of_ash.quantile(0.75) + (IQR * 3)
print('Alcalinity_of_ash outliers are values < {lowerboundary} or > {upperboundary}'.format(lowerboundary=Lower_fence, upperboundary=Upper_fence))

Alcalinity_of_ash outliers are values < 4.2999999999999997 or > 34.400000000000006

IQR = wine_df.Magnesium.quantile(0.75) - wine_df.Magnesium.quantile(0.25)
Lower_fence = wine_df.Magnesium.quantile(0.25) - (IQR * 3)
Upper_fence = wine_df.Magnesium.quantile(0.75) + (IQR * 3)
print('Magnesium outliers are values < {lowerboundary} or > {upperboundary}'.format(lowerboundary=Lower_fence, upperboundary=Upper_fence))

Magnesium outliers are values < 31.0 or > 164.0

IQR = wine_df.Proanthocyanins.quantile(0.75) - wine_df.Proanthocyanins.quantile(0.25)
Lower_fence = wine_df.Proanthocyanins.quantile(0.25) - (IQR * 3)
Upper_fence = wine_df.Proanthocyanins.quantile(0.75) + (IQR * 3)
print('Proanthocyanins outliers are values < {lowerboundary} or > {upperboundary}'.format(lowerboundary=Lower_fence, upperboundary=Upper_fence))

Proanthocyanins outliers are values < -0.8499999999999996 or > 4.05

IQR = wine_df.Color_intensity.quantile(0.75) - wine_df.Color_intensity.quantile(0.25)
Lower_fence = wine_df.Color_intensity.quantile(0.25) - (IQR * 3)
Upper_fence = wine_df.Color_intensity.quantile(0.75) + (IQR * 3)
print('Color_intensity outliers are values < {lowerboundary} or > {upperboundary}'.format(lowerboundary=Lower_fence, upperboundary=Upper_fence))

Color_intensity outliers are values < -5.7200000000000015 or > 15.14

IQR = wine_df.Hue.quantile(0.75) - wine_df.Hue.quantile(0.25)
Lower_fence = wine_df.Hue.quantile(0.25) - (IQR * 3)
Upper_fence = wine_df.Hue.quantile(0.75) + (IQR * 3)
print('Hue outliers are values < {lowerboundary} or > {upperboundary}'.format(lowerboundary=Lower_fence, upperboundary=Upper_fence))

Hue outliers are values < -0.23000000000000043 or > 2.1325000000000003

```

▼ Declare feature vector and target variable

```
x = wine_df.drop(['class'], axis=1)
Y = wine_df['class']
```

▼ Split data into separate training and test set

```
from sklearn.model_selection import train_test_split

x_train, x_test, Y_train, Y_test = train_test_split(x, Y, test_size = 0.2, random_state = 0)

x_train.shape, x_test.shape

((142, 13), (36, 13))
```

▼ Engineering outliers in numerical variables

```
def max_value(df3, variable, top):
    return np.where(df3[variable]>top, top, df3[variable])

for df3 in [x_train, x_test]:
    df3['Malicacid'] = max_value(df3, 'Malicacid', 7.5)
    df3['Alcalinity_of_ash'] = max_value(df3, 'Alcalinity_of_ash', 34.4)
    df3['Magnesium'] = max_value(df3, 'Magnesium', 164)
    df3['Proanthocyanins'] = max_value(df3, 'Proanthocyanins', 4.05)
    df3['Color_intensity'] = max_value(df3, 'Color_intensity', 15.14)
    df3['Hue'] = max_value(df3, 'Hue', 2.13)

x_train.Malicacid.max(), x_test.Malicacid.max()

(5.65, 5.8)

x_train.Alcalinity_of_ash.max(), x_test.Alcalinity_of_ash.max()

(30.0, 28.5)

x_train.Magnesium.max(), x_test.Magnesium.max()

(162, 132)

x_train.Proanthocyanins.max(), x_test.Proanthocyanins.max()

(3.58, 2.45)

x_train.Color_intensity.max(), x_test.Color_intensity.max()

(13.0, 10.8)

x_train.Hue.max(), x_test.Hue.max()

(1.71, 1.38)

x_train.describe()
```

	Alcohol	Malicacid	Ash	Alcalinity_of_ash	Magnesium	Total_phenols	Flavanoids	Nonflavanoid_phenols	Proanthocyanins
count	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000
mean	12.984859	2.372606	2.366901	19.554930	100.063380	2.258662	1.949155	0.363521	1.606056
std	0.807175	1.115360	0.269684	3.442549	14.249158	0.611691	0.975921	0.127709	0.591221
min	11.030000	0.740000	1.360000	10.600000	70.000000	1.100000	0.470000	0.130000	0.420000
25%	12.347500	1.602500	2.222500	17.250000	89.000000	1.705000	1.037500	0.270000	1.242500
50%	13.040000	1.895000	2.360000	19.500000	98.000000	2.210000	2.035000	0.340000	1.555000
75%	13.637500	3.222500	2.560000	21.500000	106.750000	2.735000	2.760000	0.450000	1.957500
max	14.750000	5.650000	3.220000	30.000000	162.000000	3.880000	3.740000	0.660000	3.580000

```
cols = x_train.columns

from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()

x_train = scaler.fit_transform(x_train)

x_test = scaler.transform(x_test)

x_train = pd.DataFrame(x_train, columns=[cols])

x_test = pd.DataFrame(x_test, columns=[cols])

x_train.describe()
```

	Alcohol	Malicacid	Ash	Alcalinity_of_ash	Magnesium	Total_phenols	Flavanoids	Nonflavanoid_phenols	Proanthocyanins
count	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000
mean	0.525500	0.332506	0.541345	0.461594	0.326776	0.416785	0.452341	0.440606	0.375334
std	0.216983	0.227161	0.144991	0.177451	0.154882	0.220033	0.298447	0.240960	0.187095
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.354167	0.175662	0.463710	0.342784	0.206522	0.217626	0.173547	0.264151	0.260285
50%	0.540323	0.235234	0.537634	0.458763	0.304348	0.399281	0.478593	0.396226	0.359177
75%	0.700941	0.505601	0.645161	0.561856	0.399457	0.588129	0.700306	0.603774	0.486551
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

Model Training

```
from sklearn.linear_model import LogisticRegression

logreg= LogisticRegression(solver='liblinear', random_state=0)

logreg.fit(x_train, Y_train)
```

▼

LogisticRegression

LogisticRegression(random_state=0, solver='liblinear')

Predict results

Predict type of classes

```
y_pred_test = logreg.predict(x_test)

y_pred_test

array([1, 3, 2, 1, 2, 2, 1, 3, 2, 2, 3, 3, 1, 2, 3, 2, 1, 1, 3, 1, 2, 1,
       1, 2, 2, 2, 2, 2, 2, 3, 1, 1, 2, 1, 1, 1])
```

✓ Prediction probability of getting 1

```
logreg.predict_proba(x_test)[: ,0]

array([0.83411707, 0.08169456, 0.33794154, 0.79491944, 0.2372458 ,
       0.237406 , 0.87096795, 0.03861458, 0.16257044, 0.05946295,
       0.1482802 , 0.03716879, 0.92972498, 0.47744117, 0.0838979 ,
       0.13261824, 0.78069736, 0.95517064, 0.0825623 , 0.84310019,
       0.47020366, 0.67960296, 0.47193266, 0.24693731, 0.08366256,
       0.16146061, 0.21250699, 0.05442331, 0.07420845, 0.07434986,
       0.83480123, 0.85379414, 0.08617887, 0.8224449 , 0.87731904,
       0.6611851 ])
```

✓ Prediction probability of getting 2

```
logreg.predict_proba(x_test)[: ,1]

array([0.1240818 , 0.06966025, 0.65228402, 0.1607158 , 0.61192177,
       0.75455152, 0.08070117, 0.14839743, 0.77629276, 0.77373406,
       0.13986198, 0.07780088, 0.03752654, 0.51475621, 0.07579536,
       0.85273997, 0.16126423, 0.02246023, 0.44592902, 0.13936754,
       0.52143812, 0.24521823, 0.43573301, 0.71879674, 0.58483752,
       0.76777135, 0.73705714, 0.8411666 , 0.73358901, 0.06019109,
       0.12898216, 0.11348826, 0.61642905, 0.06319566, 0.08703462,
       0.31648131])
```

✓ Check accuracy score

```
from sklearn.metrics import accuracy_score

print('Model accuracy score: {0:0.4f}'.format(accuracy_score(Y_test, y_pred_test)))

Model accuracy score: 0.9722
```

```
y_pred_train = logreg.predict(x_train)

y_pred_train

array([3, 2, 2, 3, 1, 1, 2, 2, 2, 1, 3, 2, 3, 1, 3, 3, 1, 3, 1, 2, 3, 3,
       2, 3, 3, 1, 3, 3, 2, 3, 3, 2, 1, 2, 2, 2, 1, 1, 2, 2, 3, 3, 2, 2,
       2, 3, 3, 1, 3, 2, 2, 2, 2, 1, 1, 2, 1, 3, 1, 3, 1, 1, 2, 1, 2,
       2, 1, 3, 2, 1, 2, 2, 2, 3, 1, 3, 3, 1, 1, 2, 3, 1, 1, 2, 2, 1, 1,
       1, 3, 2, 1, 2, 3, 1, 2, 3, 3, 1, 1, 3, 1, 3, 2, 1, 1, 2, 1, 3, 2,
       3, 1, 3, 3, 3, 1, 2, 2, 2, 2, 3, 3, 2, 2, 2, 2, 3, 3, 1, 1, 3, 2,
       2, 2, 1, 1, 1, 2, 2, 2, 1, 3])
```

```
print('Training-set accuracy score: {0:0.4f}'.format(accuracy_score(Y_train, y_pred_train)))

Training-set accuracy score: 0.9789
```

✓ Check for overfitting and underfitting

```
print('Training set score: {:.4f}'.format(logreg.score(x_train, Y_train)))

print('Test set score: {:.4f}'.format(logreg.score(x_test, Y_test)))

Training set score: 0.9789
Test set score: 0.9722
```

```
logreg100 = LogisticRegression(C=100, solver='liblinear', random_state=0)

logreg100.fit(x_train, Y_train)
```

```
▼ LogisticRegression
LogisticRegression(C=100, random_state=0, solver='liblinear')
```

```
print('Training set score: {:.4f}'.format(logreg100.score(x_train, Y_train)))
```

```
print('Test set score: {:.4f}'.format(logreg100.score(x_test, Y_test)))
```

```
Training set score: 1.0000
Test set score: 1.0000
```

```
logreg001 = LogisticRegression(C=0.01, solver='liblinear', random_state=0)
```

```
logreg001.fit(x_train, Y_train)
```

```
▼ LogisticRegression
LogisticRegression(C=0.01, random_state=0, solver='liblinear')
```

```
print('Training set score: {:.4f}'.format(logreg001.score(x_train, Y_train)))
```

```
print('Test set score: {:.4f}'.format(logreg001.score(x_test, Y_test)))
```

```
Training set score: 0.8944
Test set score: 0.8889
```

▼ Classification Report

```
from sklearn.metrics import classification_report
```

```
print(classification_report(Y_test, y_pred_test))
```

	precision	recall	f1-score	support
1	1.00	1.00	1.00	14
2	1.00	0.94	0.97	16
3	0.86	1.00	0.92	6
accuracy			0.97	36
macro avg	0.95	0.98	0.96	36
weighted avg	0.98	0.97	0.97	36

▼ Classification accuracy

```
classification_accuracy = (TP + TN) / float(TP + TN + FP + FN)
```

```
print('Classification accuracy : {0:0.4f}'.format(classification_accuracy))
```

```
Classification accuracy : 1.0000
```

▼ Precision

```
precision = TP / float(TP+FP)
```

```
print('Precision: {0:0.4f}'.format(precision))
```

```
Precision: 1.0000
```

▼ Recall

```
recall = TP / float(TP + FN)
```

```
print('Recall or Sensitivity : {0:0.4f}'.format(recall))
```

```
Recall or Sensitivity : 1.0000
```

✓ True Positive Rate

```
true_positive_rate = TP / float(TP + FN)
```

```
print('True Positive Rate : {0:0.4f}'.format(true_positive_rate))
```

```
True Positive Rate : 1.0000
```

✓ Specificity

```
specificity = TN / (TN+FP)
```

```
print('Specifity : {0:0.4f}'.format(specificity))
```

```
Specifity : 1.0000
```

✓ Adjusting the threshold level

```
y_pred_prob = logreg.predict_proba(x_test)[0:10]
```

```
y_pred_prob
```

```
array([[0.83411707, 0.1240818 , 0.04180113],
       [0.08169456, 0.06966025, 0.84864519],
       [0.33794154, 0.65228402, 0.00977444],
       [0.79491944, 0.1607158 , 0.04436476],
       [0.2372458 , 0.61192177, 0.15083243],
       [0.237406 , 0.75455152, 0.00804248],
       [0.87096795, 0.08070117, 0.04833088],
       [0.03861458, 0.14839743, 0.81298799],
       [0.16257044, 0.77629276, 0.0611368 ],
       [0.05946295, 0.77373406, 0.16680299]])
```

```
logreg.predict_proba(x_test)[0:10, 1]
```

```
array([0.1240818 , 0.06966025, 0.65228402, 0.1607158 , 0.61192177,
       0.75455152, 0.08070117, 0.14839743, 0.77629276, 0.77373406])
```

```
y_pred1 = logreg.predict_proba(x_test)[: , 1]
```

```
plt.rcParams['font.size'] = 12
```

```
plt.hist(y_pred1, bins = 10)
```


```
plt.title('Histogram of predicted probabilities of class')
```

```
plt.xlim(0,1)
```

```
plt.xlabel('Predicted probabilities of class')
```

```
plt.ylabel('Frequency')
```

```
Text(0, 0.5, 'Frequency')
```



✓ k-Fold Cross Validation

```
from sklearn.model_selection import cross_val_score

scores = cross_val_score(logreg, x_train, Y_train, cv=5, scoring = 'accuracy')

print('Cross-validation scores:{}'.format(scores))

Cross-validation scores:[0.93103448 1.          0.96428571 1.          0.92857143]

print('Average cross-validation score: {:.4f}'.format(scores.mean()))

Average cross-validation score: 0.9648
```

✓ Hyperparameter Optimization using GridSearch CV

```
from sklearn.model_selection import GridSearchCV

parameters = [{'penalty':['l1','l2']},
               {'C':[1, 10, 100, 1000]}]

grid_search = GridSearchCV(estimator = logreg,
                           param_grid = parameters,
                           scoring = 'accuracy',
                           cv = 5,
                           verbose=0)

grid_search.fit(x_train, Y_train)
```



```

/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:378: FitFailedWarning:
10 fits failed out of a total of 30.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.

```

Below are more details about the failures:

5 fits failed with the following error:

Traceback (most recent call last):

```

File "/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File "/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1160, in fit
    self._validate_params()
File "/usr/local/lib/python3.10/dist-packages/sklearn/base.py", line 600, in _validate_params
    validate_parameter_constraints(
File "/usr/local/lib/python3.10/dist-packages/sklearn/utils/_param_validation.py", line 97, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'penalty' parameter of LogisticRegression must be a str among {'none' (deprec

```

5 fits failed with the following error:

Traceback (most recent call last):

```


File "/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File "/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1160, in fit
    self._validate_params()
File "/usr/local/lib/python3.10/dist-packages/sklearn/base.py", line 600, in _validate_params
    validate_parameter_constraints(
File "/usr/local/lib/python3.10/dist-packages/sklearn/utils/_param_validation.py", line 97, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'penalty' parameter of LogisticRegression must be a str among {'none' (deprec

```

```

warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:952: UserWarning: One or more of the test scores are non-fini
warnings.warn(

```



```

print('GridSearch CV best score : {:.4f}\n\n'.format(grid_search.best_score_))

print('Parameters that give the best results :', '\n\n', (grid_search.best_params_))

print('\n\nEstimator that was chosen by the search :', '\n\n', (grid_search.best_estimator_))

GridSearch CV best score : 0.9650

Parameters that give the best results :

{'C': 10}

Estimator that was chosen by the search :

LogisticRegression(C=10, random_state=0, solver='liblinear')

print('GridSearch CV score on test set: {0:0.4f}'.format(grid_search.score(x_test, Y_test)))

GridSearch CV score on test set: 1.0000

```

✓ Results and Conclusion

The selected dependent variable which is "horsepower" could be predicted successfully using all the other independent variables in linear regression. On the other hand, the model accuracy score after using logistic regression is 97.22% which yields that the classes of the columns are predicted accurately. The classification accuracy, precision, recall, true positive rate, and specificity are all 100% which means the model achieved perfect performance on the dataset. The GridSearch CV best score also yields 96.50%. The model has learned to perfectly discriminate between the classes in the dataset without making any errors.